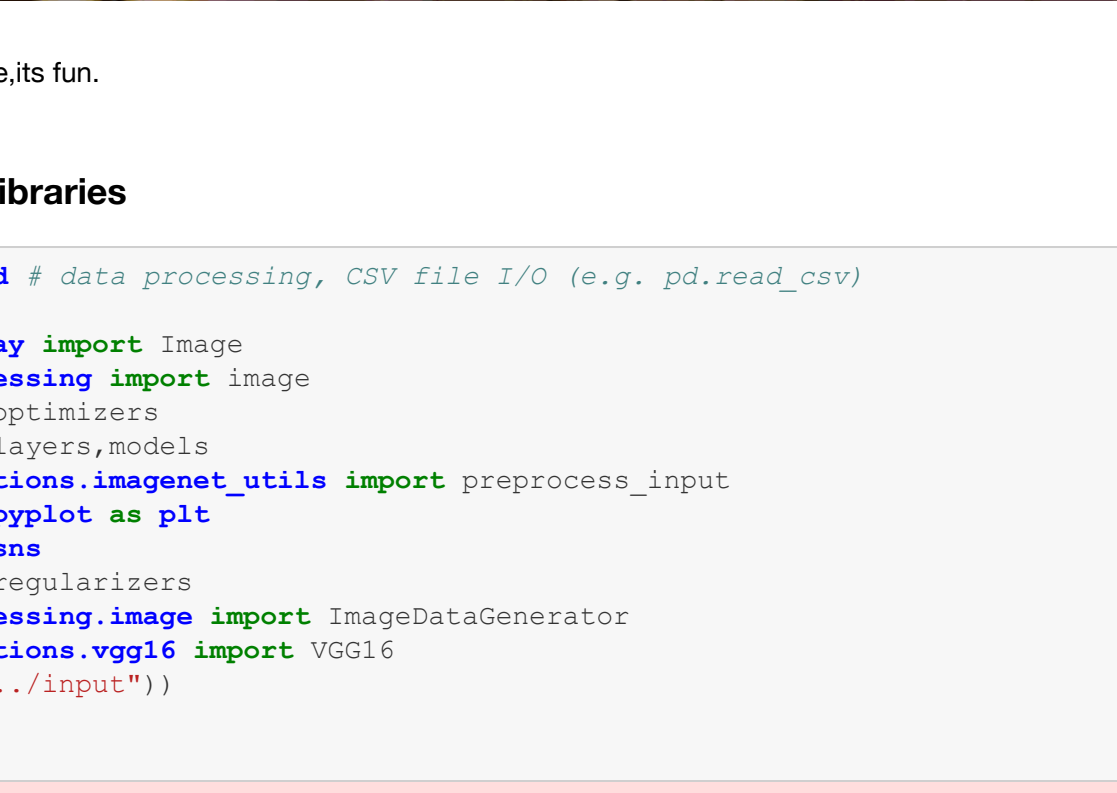


Introduction

In this kernel, we will be learning about Convolution neural networks. We will learn

- [how to prepare your data to feed to a neural network](#)
- [how convolution neural network works](#)
- [how to implement cnn in keras](#)
- [how to evaluate your results](#)
- [how to use keras pretrained network to improve your model](#)
- [how to predict and create a submission](#)

if you like my kernel, please consider upvoting



and it's no rocket science, it's fun.

Loading required libraries

```
In [1]: import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os, cv2
from IPython.display import Image
from keras.preprocessing import image
from keras.optimizers import Adam
from keras import layers, models
from keras.applications.imagenet_utils import preprocess_input
import matplotlib.pyplot as plt
import seaborn as sns
from keras import regularizers
from keras.preprocessing.image import ImageDataGenerator
from keras.applications.vgg16 import VGG16
print(os.listdir("../input"))

import numpy as np

Using TensorFlow backend.
['test', 'train', 'train.csv', 'sample_submission.csv']
```

Setting our directories

```
In [2]: train_dir = "../input/train/train"
test_dir = "../input/test/test"
train = pd.read_csv("../input/train.csv")
df_test = pd.read_csv("../input/sample_submission.csv")

The number of rows in test set is 4000
The number of rows in rows set is 4000
```

Getting a basic idea

```
In [3]: train.head(5)
train.has_cactus=train.has_cactus.astype(str)

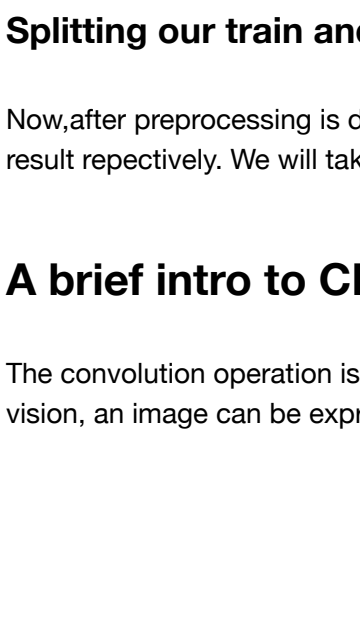
In [4]: print("out dataset has {} rows and {} columns".format(train.shape[0], train.shape[1]))
out dataset has 17500 rows and 2 columns

In [5]: train['has_cactus'].value_counts()
Out[5]: 1    13136
        0    4364
        Name: has_cactus, dtype: int64

In [6]: print("The number of rows in test set is {}".format(len(os.listdir("../input/test/test"))))
The number of rows in test set is 4000
```

Displaying an image

```
In [7]: Image(os.path.join("../input/train/train", train.iloc[0,0]), width=250, height=250)
Out[7]:
```



Data preparation

As you know, data should be processed into appropriately pre-processed floating point tensors before being fed to our network. So the steps for getting it into our network are roughly

- Read the picture files
- Decode JPEG content to RGB pixels
- Convert this into floating tensors
- Rescale pixel values (between 0 to 255) to [0,1] interval.

We will make use of ImageDataGenerator method available in keras to do all the preprocessing.

```
In [9]: datagen=ImageDataGenerator(rescale=1./255)
        batch_size=150

flow_from_dataframe Method
```

This method is useful when the images are clustered in only one folder. To put in other words, images from different classes/labels reside in only one folder. Generally, with such kind of data, some text files containing information on class and other parameters are provided. In this case, we will create a dataframe using pandas and text files provided, and create a meaningful dataframe with columns having file name (only the file names, not the path) and other classes to be used by the model. For this method, arguments to be used are:

```
dataframe value : Dataframe having meaningful data (file name, class columns are a must)
directory value : The path to the parent directory containing all images.
x_col value : which will be the name of column (in dataframe) having file names
y_col value : which will be the name of column (in dataframe) having class/label
```

```
In [9]: train_generator=datagen.flow_from_dataframe(dataframe=train[:15001], directory=train_dir, x_col='id',
        y_col='has_cactus', class_mode='binary', batch_size=batch_size,
        target_size=(150,150))

validation_generator=datagen.flow_from_dataframe(dataframe=train[15000:], directory=train_dir, x_col='id',
        y_col='has_cactus', class_mode='binary', batch_size=50,
        target_size=(150,150))

Found 15001 images belonging to 2 classes.
Found 2500 images belonging to 2 classes.
```

Splitting our train and validation dataset

Now, after preprocessing is done with our data, we will split our dataset to training and validation for training our model and validating the result respectively. We will take first 15000 images to train our data and last 2500 images to validate our model later.

A brief intro to CNN

The convolution operation is the building block of a convolutional neural network as the name suggests it. Now, in the field of computer vision, an image can be expressed as a matrix of RGB values. Therefore, let's consider the 6x6 matrix below as a part of an image:

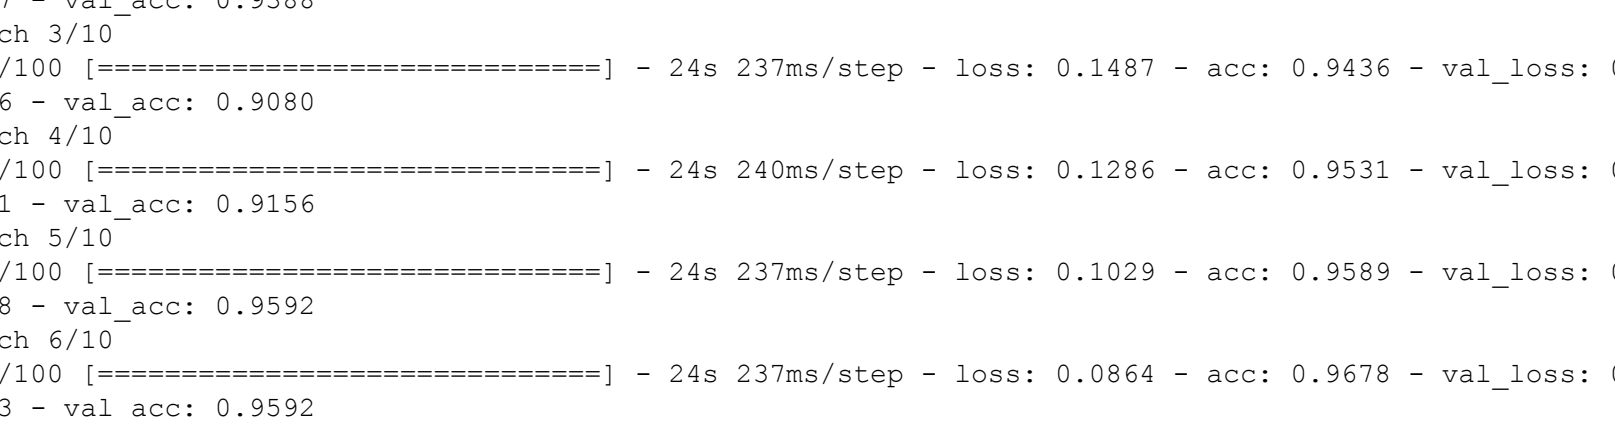
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

And the filter will be the following matrix:

1	0	-1
1	0	-1
1	0	-1

Then, the convolution involves superimposing the filter onto the image matrix, adding the product of the values from the filter and the values from the image matrix, which will generate a 4x4 convoluted layer.

This is very hard to put in words, but here is a nice animation that explains the convolution:

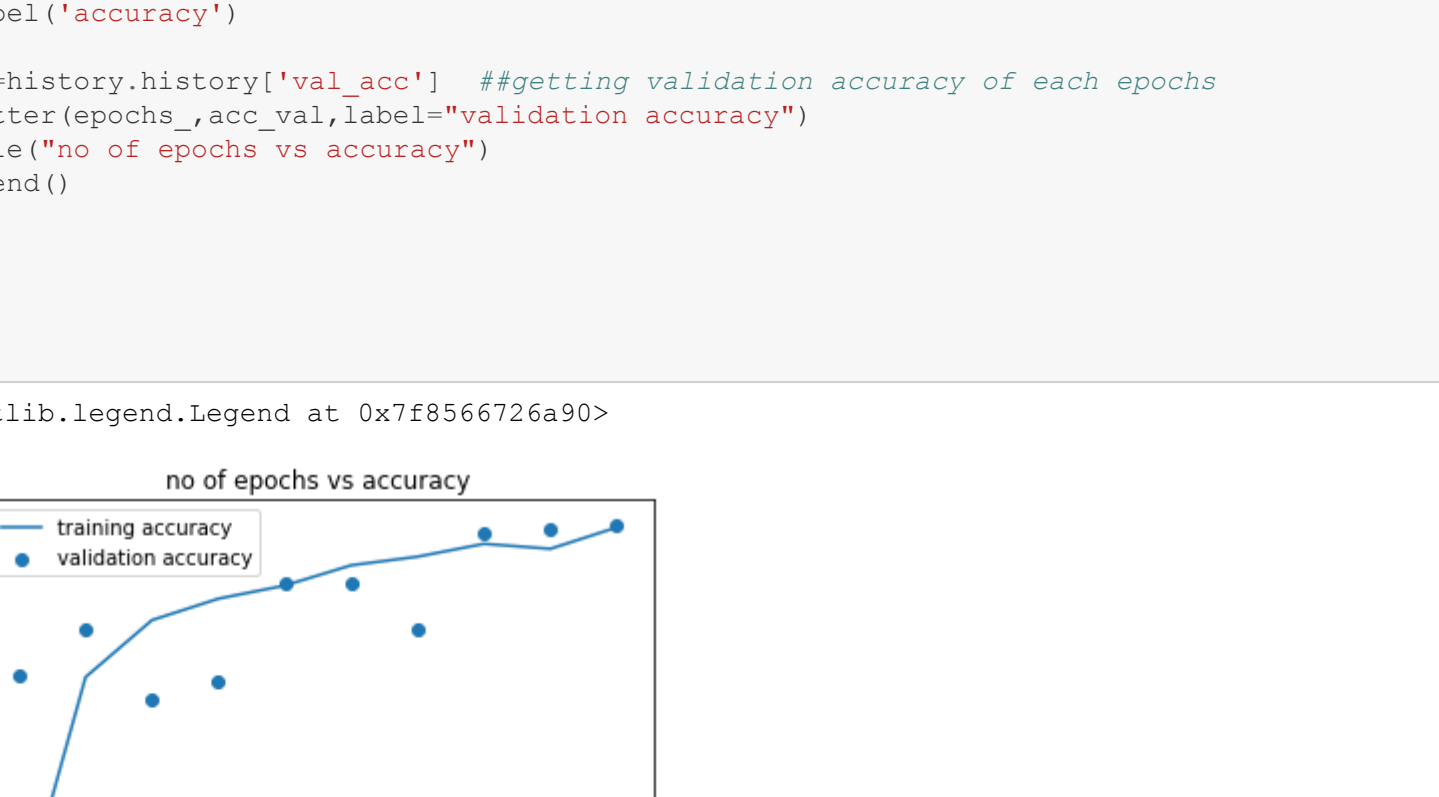


Convolutions are defined on two key parameters

- The size of patches that are extracted from input feature map, i.e. here 3x3
- The number of filters computed from convolutions

Maxpooling

Maxpooling consists of extracting features from input feature map and outputting maximum value of each channel.



Building our model

Now we will build our network. We will build our model such that it contains 5 Conv2D + Maxpooling2D stages with relu activation function.

```
In [10]: model=models.Sequential()
        model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(150,150,3)))
        model.add(layers.MaxPool2D((2,2)))
        model.add(layers.Conv2D(64, (3,3), activation='relu', input_shape=(150,150,3)))
        model.add(layers.MaxPool2D((2,2)))
        model.add(layers.Conv2D(128, (3,3), activation='relu', input_shape=(150,150,3)))
        model.add(layers.MaxPool2D((2,2)))
        model.add(layers.Conv2D(128, (3,3), activation='relu', input_shape=(150,150,3)))
        model.add(layers.MaxPool2D((2,2)))
        model.add(layers.Flatten())
        model.add(layers.Dense(512, activation='relu'))
        model.add(layers.Dense(1, activation='sigmoid'))

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
Colocations handled automatically by placer.
```

Displaying summary of our network

```
In [11]: model.summary()

Layer (type) Output Shape Param #
-----
conv2d_1 (Conv2D) (None, 148, 148, 32) 896
max_pooling2d_1 (MaxPooling2D) (None, 74, 74, 32) 0
conv2d_2 (Conv2D) (None, 72, 72, 64) 18496
max_pooling2d_2 (MaxPooling2D) (None, 36, 36, 64) 0
conv2d_3 (Conv2D) (None, 34, 34, 128) 73856
max_pooling2d_3 (MaxPooling2D) (None, 17, 17, 128) 0
conv2d_4 (Conv2D) (None, 15, 15, 128) 147584
max_pooling2d_4 (MaxPooling2D) (None, 7, 7, 128) 0
flatten_1 (Flatten) (None, 6272) 0
dense_1 (Dense) (None, 512) 3211776
-----
Total params: 3,453,121
Trainable params: 3,453,121
Non-trainable params: 0
```

Compiling our model

```
In [12]: model.compile(loss='binary_crossentropy', optimizer=optimizers.rmsprop(), metrics=['acc'])

In this step we will specify 3 important things related to our model
```

- loss: we will set our loss as binary_crossentropy since we are attacking a binary classification problem
- optimizer: we will set solver and mold our model into its most accurate possible form by fudging with the weights.
- metrics: This is the evaluation criteria that we choose to evaluate our model

```
In [12]: model.compile(loss='binary_crossentropy', optimizer=optimizers.rmsprop(), metrics=['acc'])

Fitting our model
```

Here we will use keras fit_generator() method instead of fit() method because we have used ImageDataGenerator to generate values.

```
In [13]: epochs=10
        history=model.fit_generator(train_generator, steps_per_epoch=100, epochs=10, validation_data=(validation_generator, validation_labels),
        validation_steps=50)

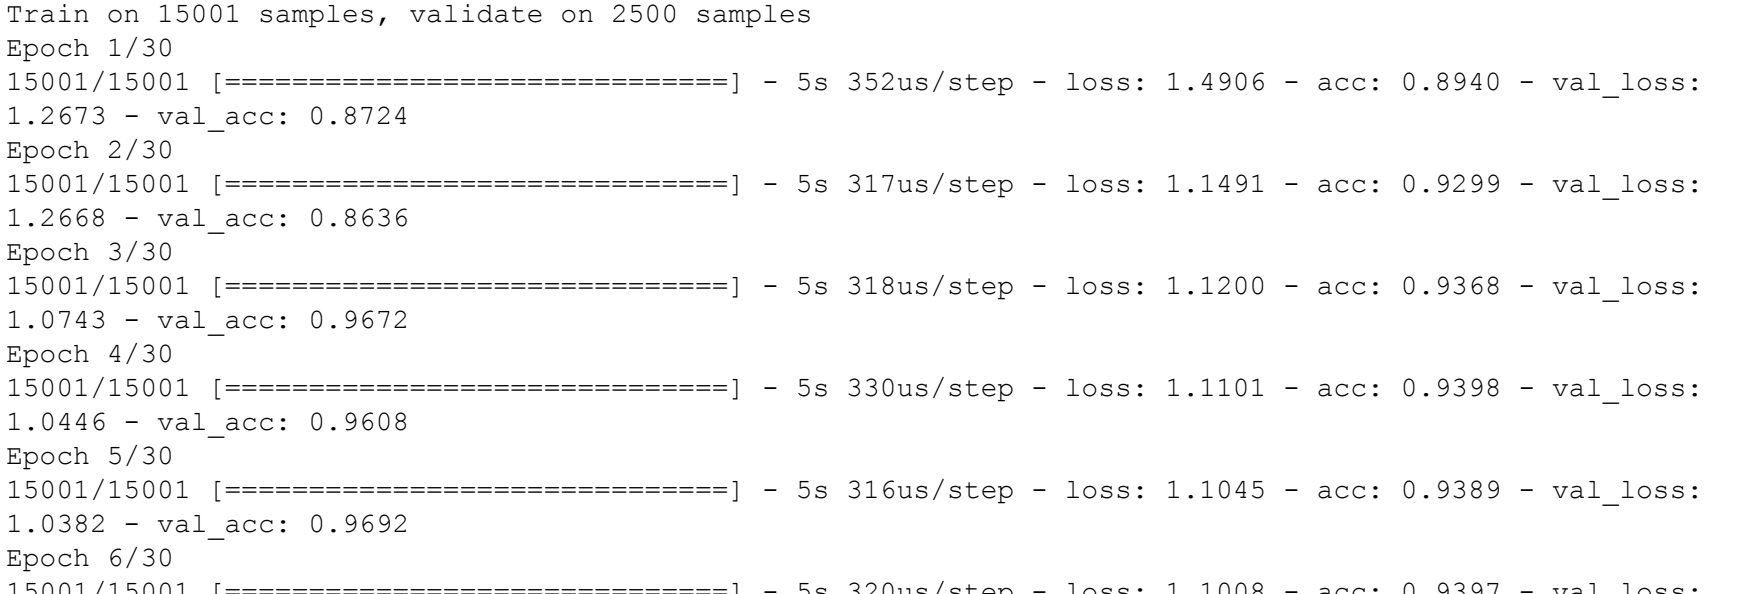
WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/ops/math_ops.py:306: tf.nn.conv2d (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
Colocations handled automatically by placer.

Epoch 1/10
100/100 [=====] - 39s 394ms/step - loss: 0.5635 - acc: 0.8113 - val_loss: 0.2326 - val_acc: 0.9188
Epoch 2/10
2846 - val_acc: 0.9080
Epoch 4/10
100/100 [=====] - 24s 240ms/step - loss: 0.1286 - acc: 0.9531 - val_loss: 0.2111 - val_acc: 0.9156
Epoch 5/10
100/100 [=====] - 24s 237ms/step - loss: 0.1029 - acc: 0.9589 - val_loss: 0.1168 - val_acc: 0.9592
Epoch 6/10
100/100 [=====] - 24s 237ms/step - loss: 0.0864 - acc: 0.9678 - val_loss: 0.1003 - val_acc: 0.9592
Epoch 7/10
100/100 [=====] - 24s 236ms/step - loss: 0.0872 - acc: 0.9716 - val_loss: 0.10247 - val_acc: 0.9388
Epoch 8/10
100/100 [=====] - 24s 239ms/step - loss: 0.0801 - acc: 0.9773 - val_loss: 0.0626 - val_acc: 0.9820
Epoch 9/10
100/100 [=====] - 24s 239ms/step - loss: 0.0571 - acc: 0.9751 - val_loss: 0.0681 - val_acc: 0.9840
Epoch 10/10
100/100 [=====] - 24s 237ms/step - loss: 0.0459 - acc: 0.9845 - val_loss: 0.0636 - val_acc: 0.9852
```

Evaluating our model


We will plot our results using matplotlib. First we will plot training and validation accuracy.

```
In [14]: acc=history.history['acc'] #getting accuracy of each epochs
        epochs=range(0, epochs)
        plt.plot(epochs, acc, label='training accuracy')
        plt.xlabel('no of epochs')
        plt.ylabel('accuracy')
```



Now we will plot training loss and validation loss of number of epochs

```
In [15]: acc=history.history['loss'] #getting loss of each epochs
        epochs=range(0, epochs)
        plt.plot(epochs, acc, label='training loss')
        plt.xlabel('No of epochs')
        plt.ylabel('loss')
```



Improving our model using VGG16

- Instantiating the VGG16 convolution base

```
In [16]: model_vg=VGG16(weights='imagenet', include_top=False)
        model_vg.summary()

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
5899238/58889256 [=====] - 1s 0us/step

Layer (type) Output Shape Param #
-----
input_1 (InputLayer) (None, None, None, 3) 0
block1_conv1 (Conv2D) (None, None, None, 64) 192
block1_conv2 (Conv2D) (None, None, None, 64) 36928
block1_pool1 (MaxPooling2D) (None, None, None, 64) 0
block2_conv1 (Conv2D) (None, None, None, 128) 73856
block2_conv2 (Conv2D) (None, None, None, 128) 147584
block2_pool1 (MaxPooling2D) (None, None, None, 128) 0
block3_conv1 (Conv2D) (None, None, None, 256) 295168
block3_conv2 (Conv2D) (None, None, None, 256) 590080
block3_conv3 (Conv2D) (None, None, None, 256) 590080
block3_pool1 (MaxPooling2D) (None, None, None, 256) 0
block4_conv1 (Conv2D) (None, None, None, 512) 1180160
block4_conv2 (Conv2D) (None, None, None, 512) 2359808
block4_conv3 (Conv2D) (None, None, None, 512) 2359808
block4_pool1 (MaxPooling2D) (None, None, None, 512) 0
block5_conv1 (Conv2D) (None, None, None, 512) 2359808
block5_conv2 (Conv2D) (None, None, None, 512) 2359808
block5_conv3 (Conv2D) (None, None, None, 512) 2359808
block5_pool1 (MaxPooling2D) (None, None, None, 512) 0
-----
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
```

Extracting features using VGG16

```
In [17]: def extract_features(directory, samples, df):

        features=np.zeros(shape=(samples, 4, 4, 512))
        labels=np.zeros(shape=samples)
        generator=datagen.flow_from_dataframe(dataframe=df, directory=directory, x_col='id',
        y_col='has_cactus', class_mode='other', batch_size=batch_size,
        target_size=(150,150))

        i=0
        for input_batch, label_batch in generator:
            feature_batch=model_vg.predict(input_batch)
            features[i*batch_size:(i+1)*batch_size]=feature_batch
            labels[i*batch_size:(i+1)*batch_size]=label_batch
            i+=1
        if i*batch_size==samples:
            break
        return(features, labels)

train.has_cactus=train.has_cactus.astype(int)
features, labels=extract_features(train_dir, 17500, train)
train_features=features[:15001]
train_labels=labels[:15001]

validation_features=features[15000:]
validation_labels=labels[15000:]

Found 17500 images.
```

Reshaping our features to feed into our dense layers

```
In [18]: #df_test.has_cactus=df_test.has_cactus.astype(str)
        test_features, test_labels=extract_features(test_dir, 4000, df_test)
        Found 4000 images.
```

```
In [19]: train_features=train_features.reshape((15001, 4*4*512))
        validation_features=validation_features.reshape((2500, 4*4*512))
        test_features=test_features.reshape((4000, 4*4*512))
```

Define a densely connected network

```
In [20]: model=models.Sequential()
        model.add(layers.Dense(212, activation='relu', kernel_regularizer=regularizers.l2_l2(.001), input_dim=(4*4*512)))
        model.add(layers.Dropout(0.2))
        model.add(layers.Dense(1, activation='sigmoid'))
```

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:34: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use 'rate' instead of 'keep_prob'. Rate should be set to 'rate = 1 - keep_prob'.

```
In [21]: model.compile(optimizer=optimizers.rmsprop(), loss='binary_crossentropy', metrics=['acc'])

In [22]: history=model.fit(train_features, train_labels, epochs=30, batch_size=15, validation_data=(validation_features, validation_labels))

Train on 15001 samples, validate on 2500 samples
15001/15001 [=====] - 5s 352us/step - loss: 1.4906 - acc: 0.8940 - val_loss: 1.2673 - val_acc: 0.8724
Epoch 2/30
15001/15001 [=====] - 5s 317us/step - loss: 1.1491 - acc: 0.9299 - val_loss: 1.0619 - val_acc: 0.9672
Epoch 3/30
15001/15001 [=====] - 5s 318us/step - loss: 1.1200 - acc: 0.9368 - val_loss: 1.0397 - val_acc: 0.9672
Epoch 4/30
15001/15001 [=====] - 5s 330us/step - loss: 1.1101 - acc: 0.9398 - val_loss: 1.0446 - val_acc: 0.9608
Epoch 5/30
15001/15001 [=====] - 5s 316us/step - loss: 1.1045 - acc: 0.9389 - val_loss: 1.0382 - val_acc: 0.9692
Epoch 6/30
15001/15001 [=====] - 5s 320us/step - loss: 1.1008 - acc: 0.9397 - val_loss: 1.0140 - val_acc: 0.9224
Epoch 7/30
15001/15001 [=====] - 5s 317us/step - loss: 1.1012 - acc: 0.9401 - val_loss: 1.1909 - val_acc: 0.8972
Epoch 8/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0972 - acc: 0.9407 - val_loss: 1.0379 - val_acc: 0.9652
Epoch 9/30
15001/15001 [=====] - 5s 317us/step - loss: 1.0961 - acc: 0.9403 - val_loss: 1.0513 - val_acc: 0.9632
Epoch 10/30
15001/15001 [=====] - 5s 314us/step - loss: 1.0897 - acc: 0.9437 - val_loss: 1.0397 - val_acc: 0.9596
Epoch 11/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0926 - acc: 0.9412 - val_loss: 1.0649 - val_acc: 0.9540
Epoch 12/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0878 - acc: 0.9423 - val_loss: 1.0963 - val_acc: 0.9420
Epoch 13/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0902 - acc: 0.9445 - val_loss: 1.0247 - val_acc: 0.9620
Epoch 14/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0956 - acc: 0.9426 - val_loss: 1.0148 - val_acc: 0.9676
Epoch 15/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0896 - acc: 0.9427 - val_loss: 1.1422 - val_acc: 0.9204
Epoch 16/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0941 - acc: 0.9416 - val_loss: 1.0180 - val_acc: 0.9636
Epoch 17/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0938 - acc: 0.9399 - val_loss: 1.0270 - val_acc: 0.9244
Epoch 18/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0907 - acc: 0.9425 - val_loss: 1.0147 - val_acc: 0.9720
Epoch 19/30
15001/15001 [=====] - 5s 318us/step - loss: 1.0908 - acc: 0.9407 - val_loss: 1.0156 - val_acc: 0.9704
Epoch 20/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0907 - acc: 0.9441 - val_loss: 1.0397 - val_acc: 0.9604
Epoch 21/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0965 - acc: 0.9395 - val_loss: 1.0619 - val_acc: 0.9604
Epoch 22/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0938 - acc: 0.9399 - val_loss: 1.0146 - val_acc: 0.9732
Epoch 23/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0966 - acc: 0.9417 - val_loss: 1.0270 - val_acc: 0.9652
Epoch 24/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0896 - acc: 0.9427 - val_loss: 1.1422 - val_acc: 0.9204
Epoch 25/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0941 - acc: 0.9416 - val_loss: 1.0180 - val_acc: 0.9636
Epoch 26/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0938 - acc: 0.9399 - val_loss: 1.0270 - val_acc: 0.9676
Epoch 27/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0876 - acc: 0.9452 - val_loss: 1.0180 - val_acc: 0.9636
Epoch 28/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0897 - acc: 0.9441 - val_loss: 1.0140 - val_acc: 0.9244
Epoch 29/30
15001/15001 [=====] - 5s 315us/step - loss: 1.0879 - acc: 0.9435 - val_loss: 1.0241 - val_acc: 0.9636
Epoch 30/30
15001/15001 [=====] - 5s 316us/step - loss: 1.0930 - acc: 0.9423 - val_loss: 1.2741 - val_acc: 0.8884
```

making prediction

```
In [23]: y_pre=model.predict_proba(test_features)
```

Making a submission

```
In [24]: df=pd.DataFrame({'id':df_test['id'] })
        df['has_cactus']=y_pre
        df.to_csv("submission.csv", index=False)
```

If you guys like my kernel, please do consider upvoting it.

```
In [25]:
```